

+

Machine Learning and Data Mining

Dimensionality Reduction; PCA & SVD

Prof. Alexander Ihler





Motivation

- High-dimensional data
 - Images of faces
 - Text from articles
 - All S&P 500 stocks
- Can we describe them in a “simpler” way?
- Ex: S&P 500 – vector of 500 (change in) values per day
 - But, lots of structure
 - Some elements tend to “change together”
 - Maybe we only need a few values to approximate it?
 - “Tech stocks up 2x, manufacturing up 1.5x, ...” ?
- How can we access that structure?

Dimensionality reduction

- Ex: data with two real values $[x_1, x_2]$
- We'd like to describe each point using only one value $[z_1]$
- We'll communicate a "model" to convert: $[x_1, x_2] \sim f(z_1)$
- Ex: linear function $f(z)$: $[x_1, x_2] = z * \underline{v} = z * [v_1, v_2]$
- \underline{v} is the same for all data points (communicate once)
- z tells us the closest point on v to the original point $[x_1, x_2]$




Principal Components Analysis

- What is the vector that would most closely reconstruct X?

$$\min_{a,v} \sum_i (x^{(i)} - a^{(i)}v)^2$$

- Given v : $a^{(i)}$ is the projection of each point $x^{(i)}$ onto v
- v chosen to minimize the residual variance
- Equivalently, v is the direction of maximum variance
- Extensions: best two dimensions: $x_i = a_i*v + b_i*w + m$



Geometry of the Gaussian

$$\Delta^2 = (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})$$

Oval shows constant Δ^2 value...

$$\boldsymbol{\Sigma} = \mathbf{U} \boldsymbol{\Lambda} \mathbf{U}^T$$


Write $\boldsymbol{\Sigma}$ in terms of eigenvectors...

$$\boldsymbol{\Sigma}^{-1} = \sum_{i=1}^D \frac{1}{\lambda_i} \mathbf{u}_i \mathbf{u}_i^T$$

Then...

$$\Delta^2 = \sum_{i=1}^D \frac{y_i^2}{\lambda_i}$$

$$y_i = \mathbf{u}_i^T (\mathbf{x} - \boldsymbol{\mu})$$



PCA representation

- Subtract data mean from each point
- (Typically) Scale each dimension by its variance
 - Helps pay less attention to magnitude of the variable
- Compute covariance matrix, $S = 1/n \sum (x_i - m)' (x_i - m)$
- Compute the k largest eigenvectors of S

$$S = V D V^T$$

[illegible]

Singular Value Decomposition

- Alternative method to calculate (still subtract mean 1st)
- Decompose $X = U S V^T$
 - Orthogonal: $X^T X = V S S V^T = V D V^T$
 - $X X^T = U S S U^T = U D U^T$
- $U \cdot S$ matrix provides coefficients
 - Example $x_i = U_{i,1} S_{11} v_1 + U_{i,2} S_{22} v_2 + \dots$
- Gives the least-squares approximation to X of this form

$$\boxed{\begin{matrix} \mathbf{X} \\ \mathbf{N \times D} \end{matrix}} \approx \boxed{\begin{matrix} \mathbf{U} \\ \mathbf{N \times K} \end{matrix}} \boxed{\begin{matrix} \mathbf{S} \\ \mathbf{K \times K} \end{matrix}} \boxed{\begin{matrix} \mathbf{V^T} \\ \mathbf{K \times D} \end{matrix}}$$


“Eigen-faces”

- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
 - 24x24 images of faces = 576 dimensional measurements



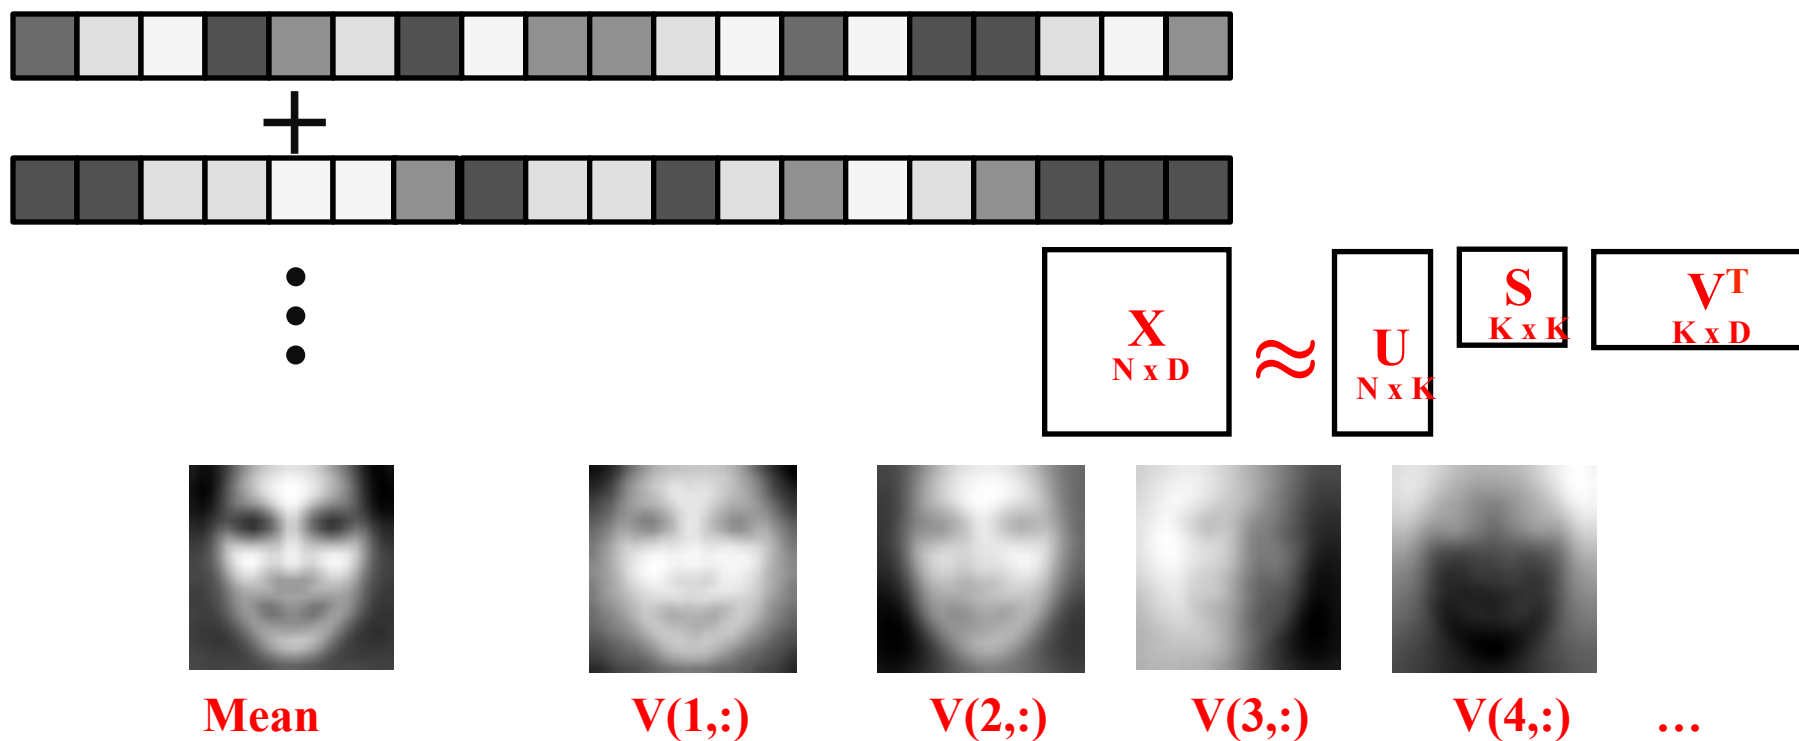
•
•
•

•
•
•



“Eigen-faces”

- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
 - 24x24 images of faces = 576 dimensional measurements
 - Take first K PCA components




“Eigen-faces”

- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
 - 24x24 images of faces = 576 dimensional measurements
 - Take first K PCA components




Mean



Dir 1



Dir 2




Dir 3




Dir 4



...



K=4



K=50



K=4



K=50

Text representations

- “Bag of words”
 - Remember word counts but not order
- Example:

Rain and chilly weather didn't keep thousands of parade-goers from camping out Friday night for the 111th Tournament of Roses.

Spirits were high among the street party crowd as they set up for curbside seats for today's parade.

“I want to party all night,” said Tyne Gaudielle, 15, of Glendale, who spent the last night of the year along Colorado Boulevard with a group of friends.

Whether they came for the partying or the parade, campers were in for a long night. Rain continued into the evening and temperatures were expected to dip down into the low 40s.

Text representations

- “Bag of words”
 - Remember word counts but not order
- Example:

Rain and chilly weather didn't keep thousands of parade-goers from camping out Friday night of Roses.

Spirits were high among the street party crowd for curbside seats for today's parade.

“I want to party all night,” said Tyne Gaudin, Glendale, who spent the last night of the year on 111th Boulevard with a group of friends.

Whether they came for the partying or the tournament in for a long night. Rain continued into the night as temperatures were expected to dip down in

example1/20000101.0015.txt

rain
chilly
weather
didn
keep
thousands
parade-goers
camping
out
friday
night
111th
tournament
roses
spirits

Text representations


- “Bag of words”
 - Remember word counts but not order
- Example:

VOCABULARY:

	DOC #	WORD #	COUNT
0001 ability	1	29	1
0002 able	1	56	1
0003 accept	1	127	1
0004 accepted	1	166	1
0005 according	1	176	1
0006 account	1	187	1
0007 accounts	1	192	1
0008 accused	1	198	2
0009 act	1	356	1
0010 acting	1	374	1
0011 action	1	381	2
0012 active	...		
....			


Latent Semantic Indexing (LSI)

- PCA for text data
- Create a giant matrix of words in docs
 - “Word j appears” = feature x_j
 - “in document i ” = data example i
- Huge matrix (mostly zeros)
 - Typically normalize by e.g. sum over j to control for short docs
 - Typically don't subtract mean or normalize by variance
 - Might transform counts in some way (log, etc)
- PCA on this matrix provides a new representation
 - Document comparison
 - Fuzzy search (“concept” instead of “word” matching)



Matrices are big, but data is sparse

- Typical example:
 - Number of docs, $D \sim 10^6$
 - Number of unique words in vocab, $W \sim 10^5$
 - FULL Storage required $\sim 10^{11}$
 - Sparse Storage required $\sim 10^9$
- $D \times W$ matrix (# docs x # words)
 - Looks dense, but that's just plotting
 - Each entry is non-negative
 - Typically integer / count data



Latent Semantic Indexing (LSI)

- What do the principal components look like?

PRINCIPAL COMPONENT 1

0.135 genetic
0.134 gene
0.131 snp
0.129 disease
0.126 genome_wide
0.117 cell
0.110 variant
0.109 risk
0.098 population
0.097 analysis
0.094 expression
0.093 gene_expression
0.092 gwas
0.089 control
0.088 human
0.086 cancer

Latent Semantic Indexing (LSI)

- What do the principal components look like?


PRINCIPAL COMPONENT 1

0.135 genetic
0.134 gene
0.131 snp
0.129 disease
0.126 genome_wide
0.117 cell
0.110 variant
0.109 risk
0.098 population
0.097 analysis
0.094 expression
0.093 gene_expression
0.092 gwas
0.089 control
0.088 human
0.086 cancer

PRINCIPAL COMPONENT 2

0.247 snp
-0.196 cell
0.187 variant
0.181 risk
0.180 gwas
0.162 population
0.162 genome_wide
0.155 genetic
0.130 loci
-0.116 mir
-0.116 expression
0.113 allele
0.108 schizophrenia
0.107 disease
-0.103 mirnas
-0.099 protein

Q: But what
does -0.196 cell
mean?



From Y. Koren
of BellKor team

Collaborative Filtering (Netflix)

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	

$$\begin{matrix} \boxed{\mathbf{X}} \\ \mathbf{N \times D} \end{matrix} \approx \begin{matrix} \boxed{\mathbf{U}} \\ \mathbf{N \times K} \end{matrix} \begin{matrix} \boxed{\mathbf{S}} \\ \mathbf{K \times K} \end{matrix} \begin{matrix} \boxed{\mathbf{V}^T} \\ \mathbf{K \times D} \end{matrix}$$

Latent Space Models

Model ratings matrix as
“user” and “movie”
positions

Infer values from known
ratings

	users											
items	1		3			5			5		4	
			5	4			4			2	1	3
	2	4		1	2		3		4	3	5	
		2	4		5			4			2	
			4	3	4	2					2	5
	1		3		3			2			4	

~

Extrapolate to unranked

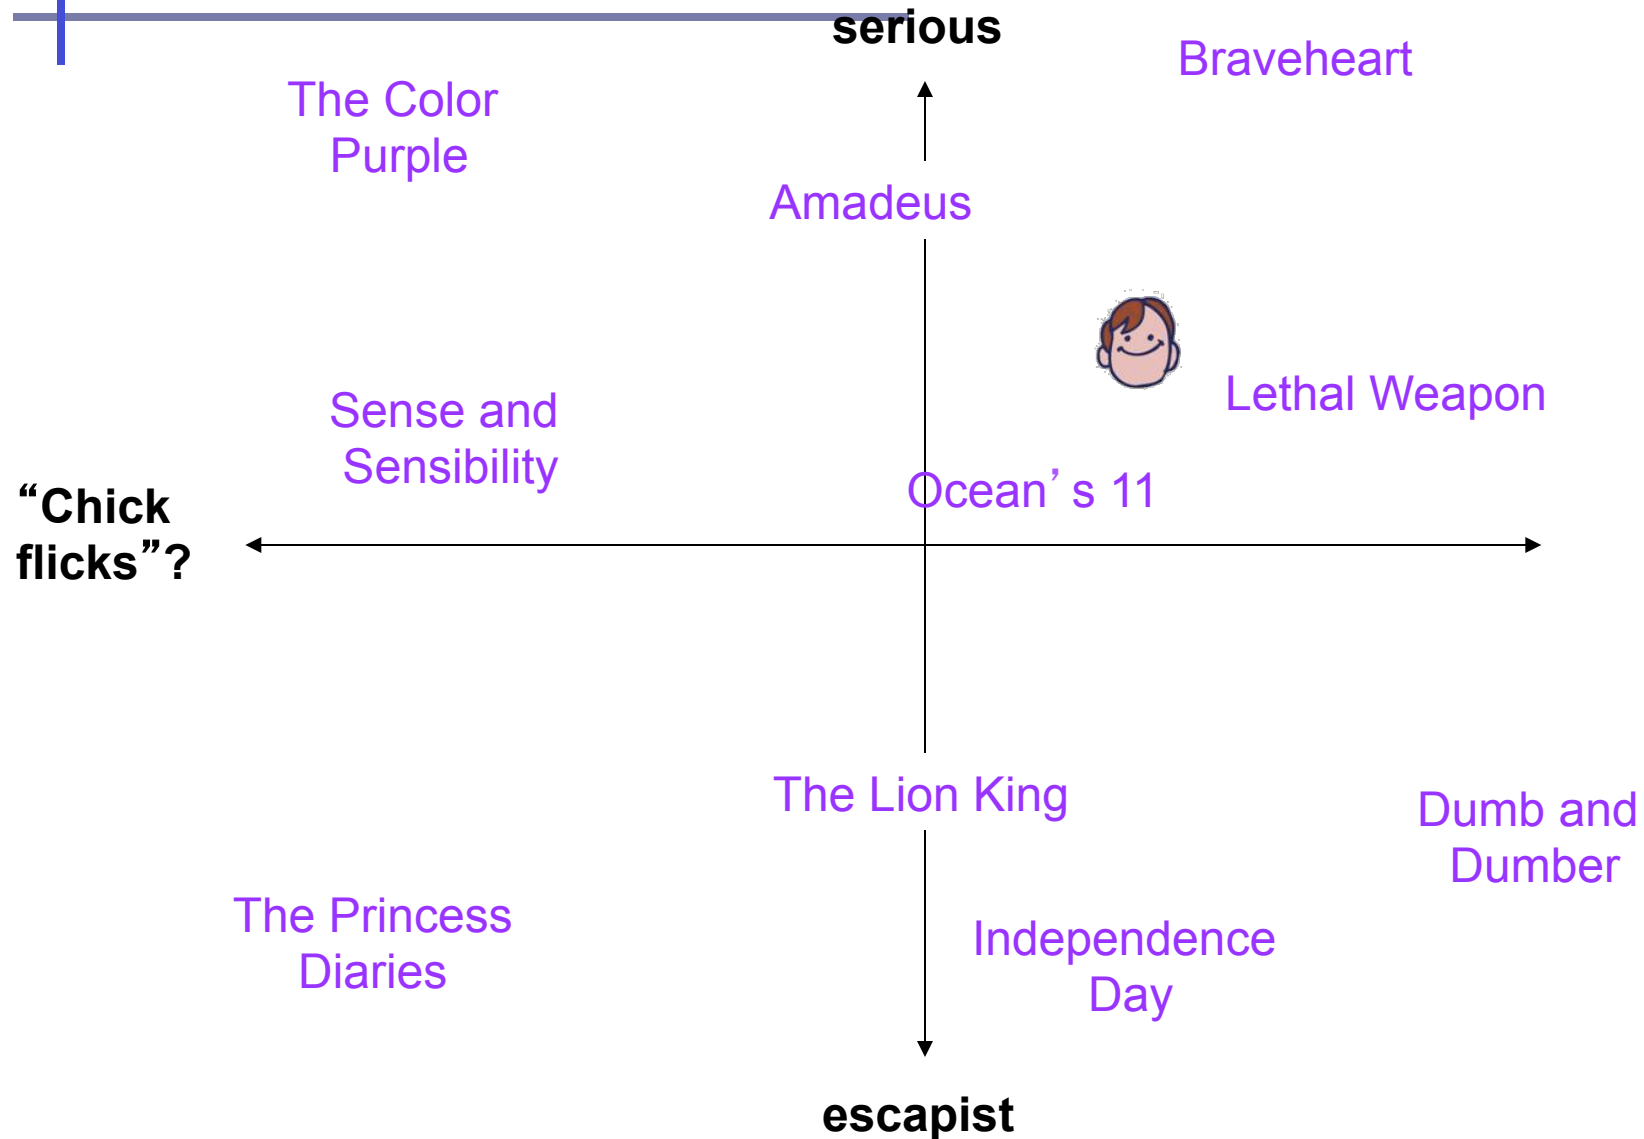
	users											
items	.1	-.4	.2									
	-.5	.6	.5									
	-.2	.3	.5									
	1.1	2.1	.3									
	-.7	2.1	-2									
	-1	.7	.3									

•

1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

From Y. Koren
of BellKor team

Latent Space Models



Some SVD dimensions

See timelydevelopment.com

Dimension 1

Offbeat / Dark-Comedy

Lost in Translation
The Royal Tenenbaums
Dogville
Eternal Sunshine of the Spotless Mind
Punch-Drunk Love

Mass-Market / 'Beniffer' Movies

Pearl Harbor
Armageddon
The Wedding Planner
Coyote Ugly
Miss Congeniality

Dimension 2

Good

VeggieTales: Bible Heroes: Lions
The Best of Friends: Season 3
Felicity: Season 2
Friends: Season 4
Friends: Season 5

Twisted

The Saddest Music in the World
Wake Up
I Heart Huckabees
Freddy Got Fingered
House of 1

Dimension 3

What a 10 year old boy would watch

Dragon Ball Z: Vol. 17: Super Saiyan
Battle Athletes Victory: Vol. 4: Spaceward Ho!
Battle Athletes Victory: Vol. 5: No Looking Back
Battle Athletes Victory: Vol. 7: The Last Dance
Battle Athletes Victory: Vol. 2: Doubt and Conflict

What a liberal woman would watch

Fahrenheit 9/11
The Hours
Going Upriver: The Long War of John Kerry
Sex and the City: Season 2
Bowling for Columbine

Latent space models

- Latent representation encodes some “meaning”
- What kind of movie is this? What movies is it similar to?
- Matrix is full of missing data
 - Hard to take SVD directly
 - Typically solve using gradient descent
 - Easy algorithm (see Netflix challenge forum)

```
%For user u, movie m, find kth eigenvector & coefficient by iterating:  
err = ( rating – predict(m,u) );           % predict vector-vector product  
    % predict(m,u) = U(m,:) * V(:,u)  
Vku = V(k,u); Umk=U(m,k);                 % save before changing  
U(m,k) = U(m,k) + alpha*err*Vku;          % Update our matrices  
V(k,u) = V(k,u) + alpha*err*Umk;          % (compare to least-squares gradient)
```

Summary

- Dimensionality reduction
 - Representation: basis vectors & coefficients
- Linear decomposition
 - PCA / eigendecomposition
 - Singular value decomposition
- Examples and data sets
 - Face images
 - Text documents (latent semantic indexing)
 - Movie ratings